

Holistic 3D Human and Scene Mesh Estimation from Single View Images

Zhenzhen Weng, Serena Yeung
Stanford University
{zwweng, syyeung}@stanford.edu

Abstract

The 3D world limits the human body pose and the human body pose conveys information about the surrounding objects. Indeed, from a single image of a person placed in an indoor scene, we as humans are adept at resolving ambiguities of the physical laws and prior perception of the plausible object and human poses. However, few computer vision models fully leverage this fact. In this work, we propose a holistically trainable model that perceives the 3D scene from a single RGB image, estimates the camera pose and the room layout, and reconstructs both human body and object meshes. By imposing a set of comprehensive and sophisticated losses on all aspects of the estimations, we show that our model outperforms existing human body mesh methods and indoor scene reconstruction methods. To the best of our knowledge, this is the first model that outputs both object and human predictions at the mesh level, and performs joint optimization on the scene and human poses.

1. Introduction

Holistic scene perception is key to our human ability to accurately interpret and interact with the 3D world. The human visual system naturally integrates context from actors, objects, and scene layout to infer realistic, robust estimations of the scene. Suppose a human is partially included in an image because they are positioned behind a desk. We can still effortlessly extract rich information from the static scene to resolve ambiguities due to the occlusion. Likewise, the appearance of humans also provides useful information about scenes, such as the ground plane and depth of surrounding objects. Humans and objects in scenes jointly manifest spatial occupancies that constrain their relative positions. For computer vision systems to achieve high accuracy in recognizing and interpreting complex scenes, it is therefore important to develop approaches for holistic scene perception and reasoning.

In recent years, holistic scene understanding from single view images has gained increasing interest from com-

puter vision researchers. [35] [14] proposed methods for joint reasoning over inanimate scenes, and recovered room layout and 3D object bounding boxes using consistency losses such as a constraint for objects to be enclosed within the room bounding box. [4] additionally discouraged intersection between object bounding box estimations, and was the first model to bring 3D human pose estimation into the holistic scene understanding problem. It incorporated human-object interaction priors to reason about approximate relations between humans and objects. However all of these works still operate at the relatively coarser level of bounding boxes and joint key points, and are therefore limited in their ability to use precise shapes, surfaces, and physical occupancies to design holistic scene constraints and improve estimation accuracy.

In this work we propose the first single-view, holistic scene understanding method that jointly optimizes over all aspects of 3D human pose, objects, and room layout at the mesh level, to produce state-of-the-art mesh estimations of the scene. Our approach builds on recent advances in mesh prediction. [3] [7] [30] proposed methods for reconstructing the individual object meshes with varying topological

structures. [28] builds on [30] and proposed the first holistic 3D scene understanding method with mesh reconstruction at the instance level, however they did not consider humans. Recently, [11] introduced a method for 3D mesh-based human pose estimation, that utilizes physical occupancy information of the static scene to discourage body penetration into the scene. However, [11] requires the ground truth 3D scans of the scene, and does not perform joint human and scene estimation.

Given a single RGB image, our method simultaneously reconstructs the human body mesh and multiple aspects of the scene – 3D object meshes and bounding boxes, room layout, and camera pose – all in 3D (Figure 1). Our approach outputs the SMPL-X (SMPL eXpressive) [31] human mesh model, which fully parameterizes the 3D surface of the human body. It also leverages a variant of the Topology Modification Network (TMN) [5], proposed in [28], as the base model for static object mesh and scene reconstruction. Importantly, we introduce a joint optimization process that incorporates a comprehensive set of physical constraints and priors including 2D/3D reprojection constraints, object-object mesh constraints, object-human mesh constraints, and object/human - room layout constraints, to obtain robust, physically plausible predictions. We perform experimental evaluation on the PiGraphs [34] and PROX [11] datasets and demonstrate that our model outperforms state-of-the-art methods on either 3D scene understanding or 3D human pose estimation.

In summary, our contributions are the following:

- We propose a holistic trainable model for jointly reconstructing 3D human body meshes and static scene elements (3D object meshes and bounding boxes, room layout, and camera pose) from monocular RGB images. To the best of our knowledge, we are the first to jointly estimate this rich scene understanding at the mesh level.
- Our model does not require any ground truth annotations of the 3D scene or the human poses, and can be directly used on any indoor dataset to produce high quality mesh reconstructions.
- Through our joint optimization process that incorporates a comprehensive set of physical constraints and priors, we show that our model outperforms prior state-of-the-art methods on either 3D scene understanding or 3D human pose estimation, on the PiGraphs and PROX Quantitative datasets.

2. Related Work

Single View 3D Human Pose Estimation. Previous 3D pose estimation methods from single view RGB images can be divided into two types: (i) directly learning 3D human keypoints from 2D image features [39], and (ii) 2D pose estimation with subsequent separate lifting of the 2D coor-

dinates to 3D via deep neural networks [33] [24].

Although these methods showed impressive results in the wild images with relatively clean backgrounds, estimating 3D poses with cluttered background and partial occlusions is still very challenging. Recent works in human body models [22] [31] and single view body mesh reconstruction methods [2] [19] have pushed the richness of body details available for reasoning, and provide opportunities for bringing novel constraints to the training stage. Recently, [11] proposed the first 3D human body mesh reconstruction method that takes the static scene into consideration; however they rely on ground truth 3-D scene scans. Our work builds on these directions and is the first to leverage mesh representations of both human and scene in performing holistic estimation of 3D human body and scene meshes jointly.

Holistic Scene Understanding. The 3D holistic scene understanding problem, in particular 3D scene reconstruction from single view images, has received increasing attention over the past few years. While most of these works have focused on coarser bounding boxes and keypoints as opposed to meshes, methods have differed in model outputs and constraint formulations [4] [28] [4]. Works such as [14] have focused on the static scene, [14] proposed an end-to-end model that learns the 3D room layout, camera pose and 3D object bounding boxes. Drawing insight from the camera projection process and physical commonsense, [14] encourages projected 3D bounding boxes to be close to their 2D locations on the image plane, and forces object bounding boxes to be within the room layout bounding box.

Some works have attempted to incorporate scene/object information in human pose estimation [43] [11] [26] [44] and/or vice-versa [8], [44] relies on mesh exemplars with annotated contact points, and does not perform full layout/scene reconstruction. [26] uses a database of “scenelets” and works with human skeletons. [11] utilizes ground truth scene scans. In contrast to these, we consider the more challenging setting of directly estimating scene and human meshes (in general indoor settings), whereas joint mesh estimation is beyond the scope of these works. [4] jointly tackles two tasks from a single-view image: (i) 3D estimations of object bounding boxes, camera pose, and room layout; and (ii) 3D human keypoints estimation. They used an energy-based inference optimization process that refines direct 3D outputs by jointly reasoning across aspects of the objects and human keypoints. However, their constraint formulations based on 3D bounding boxes and human keypoints are still lacking in precision. Additionally, energy-based models have the disadvantage of an expensive inference step compared to feed-forward models, and

are represented as $c \in \mathbb{R}^2$. Our representation for the camera pose, room layout, and 3D object bounding boxes and meshes in a scene follows the notation used in [14] [28]. The camera pose is a 3×3 rotation matrix defined by the pitch and roll angles of the camera system relative to the world system. In the world system, an object bounding box is represented by a 3D bounding box $X \in \mathbb{R}^{(8,3)}$, which can be determined from its 3D center $C \in \mathbb{R}^3$, spatial size $s \in \mathbb{R}^3$, and orientation angle $\theta \in [-\pi, \pi]$. The suboid room layout is also represented by a 3D box $X^r \in \mathbb{R}^{(8,3)}$, and is parameterized in the same manner as an object bounding box. The triangular mesh for object i in the image is represented by its vertices and faces $M_i = (V_i, F_i)$, where $V_i \in \mathbb{R}^{(N_i,3)}$, N_i is the number of vertices and F_i defines the triangular faces of the mesh. M_i is normalized to fit in a unit cube, and the vertices of the mesh can be converted to the 3D camera coordinate system by translation and rotation as specified by the 3D bounding box parameters.

Human Body Model. We represent the human body using SMPL-X (SMPL eXpressive) [31], a generative model that captures how the human body shape varies across a human population, learned from a corpus of registered 3D body, face and hand scans of people of different sizes, genders and nationalities in various poses. SMPL-X extends the SMPL model [22] with fully articulated hands and an expressive face. It is essentially a differentiable function parameterized by shape β_θ , pose θ_θ , facial expressions ψ and translation γ of the body. The output of SMPL-X is a 3D triangular mesh $M_b = (V_b, F_b)$ that contains $N_b = 10475$ vertices $V_b \in \mathbb{R}^{(N_b,3)}$ and triangular faces F_b .

3.2. Model Architecture

Body Model.

Since the SMPL-X [31] body model is a fully differentiable function, we simply compute the body loss terms (Section 3.3) that are formulated in terms of the vertices and faces of the output human body mesh, and back-propagate the SMPL-X model to find the optimal set of parameters such as the shape and pose of the human body. As in [31], the parameters of the SMPL-X model are regularized with a set of body priors including a VAE-based body pose prior, and L_2 priors on hand pose, facial pose, body shape and facial expressions, penalizing deviation from the neutral state.

Scene Models. We use three sub-modules to predict 3D object boxes, camera pose and 3D room layout, and 3D object meshes in the scene, respectively. Specifically, we adopt the Object Detection Network (ODN), Layout Estimation Network (LEN), and Mesh Generation Network (MGN) from [28]. For 3D object box prediction, the ODN first takes 2D detections of a Faster R-CNN model trained on LVIS [10], extracts appearance features in an object-wise fashion using ResNet-34 [12], and encodes the relative position and size between 2D object boxes into geometry fea-

tures using the method in [13]. For each target object, an “attention sum” is then computed using relational features to other objects [13]. Finally, each set of box parameters is regressed using a two-layer MLP. The LEN consists of a ResNet-34 feature extractor and two separate branches with fully-connected layers, one for predicting the camera pose and the other for predicting the 3D room bounding box attributes. Finally, for 3D object mesh prediction, the MGN takes a 2D detection of an object as input and uses ResNet-18 to extract 2D appearance features. Then, the image features concatenated with the one-hot LVIS [10] object category encoding are fed into the decoder of AtlasNet [9], which performs mesh deformation from a template sphere mesh. An edge classifier is trained to remove redundant edges from the deformed mesh and a boundary refinement module [30] is used to refine the smoothness of body shape, penalizing deviations from the neutral state. E_θ is a VAE-based body pose prior called VPose introduced in [31]. E_θ is a prior regularizing extended bending only for elbows and knees. The terms $E_x, E_\theta, E_\psi, E_\gamma, E_\beta$ are as described in [31]. $E_{T_{\text{joint}}}$ is a penetration penalty for self-penetrations (e.g. hand intersecting knee). The X ’s are the weights for the terms.

Our formulation is closest to that in [11], which performs human mesh estimation and was built upon [31] with the addition of scene contact (E_c) and penetration (E_p) terms by assuming access to ground truth scene scans. There are several differences between their full loss function and our

formulation in Fig. 1. First, we do not include any depth related terms, because we wish to perform estimation using solely RGB images whereas [11] propose model variants leveraging RGB-D information. Second, since we are performing joint estimation of the 3D scene from a monocular RGB image, we are not yet able to reason on scene contact or penetration after only human mesh estimation. So we include only a body self-penetration term in Fig. 1, which is computed following the approach in [1] [31] [41], and will consider human-scene constraints instead during our global optimization stage.

Within-scene losses. In Stage I of our approach, we also utilize within-scene constraints to generate an initial static scene estimation. Specifically, we design two within-scene constraints, one for encouraging 2D/3D consistency of the predicted object bounding boxes and the other one for penalizing the collision between the object meshes.

For the first constraint, we utilize the fact that based on the camera projection model, if we project predicted 3D bounding boxes onto the 2D image plane, the projected corners should be close to the 2D bounding box corners. This constraint therefore optimizes both camera pose and 3D bounding boxes. [28] imposes a similar loss where they penalize the deviation of the 2D projections of predicted 3D bounding box corners from ground truth 3D bounding box corners for both object bounding boxes and the room bounding box. However, since our model does not rely on any ground truth annotations in our described optimization process, we propose to use our detected 2D bounding boxes as a pseudo ground truth. We show the effectiveness of this loss term in Section 4. The formal definition of this term can be written as

$$\mathcal{L}_{\text{joint}}^j = \frac{1}{n_{\text{obj}}} \sum_{i=1}^{n_{\text{obj}}} \text{SmoothL}_2(f(X_i, c_i, \theta_i), h_i) \quad (2)$$

where c_i, θ_i are the size, centroid and orientation of the object i , h_i is the 2D bounding box estimate for object i , and f is a differentiable projection function that projects the corners of a 3D bounding box to a 2D image plane. Like [28], we use a smooth L_1 loss function comprised of a squared term if the absolute element-wise error falls below a threshold and an L_1 term otherwise.

Our second constraint is a loss term that penalize the collision between reconstructed object meshes. Although some pose estimation works [11] [18] have incorporated body collision losses, prior works in scene understanding have not explored this loss, because they either did not have the object shape information necessary to calculate the precise collision [18] [4], or did not take advantage of the object shape information that comes with the meshes [28]. We notice that inter-object collision is common in the output of human mesh estimation and was built upon [31] with the addition of scene contact (E_c) and penetration (E_p) terms by assuming access to ground truth scene scans. There are several differences between their full loss function and our

formulation in Fig. 1. First, we do not include any depth related terms, because we wish to perform estimation using solely RGB images whereas [11] propose model variants leveraging RGB-D information. Second, since we are performing joint estimation of the 3D scene from a monocular RGB image, we are not yet able to reason on scene contact or penetration after only human mesh estimation. So we include only a body self-penetration term in Fig. 1, which is computed following the approach in [1] [31] [41], and will consider human-scene constraints instead during our global optimization stage.

Within-scene losses. In Stage I of our approach, we also utilize within-scene constraints to generate an initial static scene estimation. Specifically, we design two within-scene constraints, one for encouraging 2D/3D consistency of the predicted object bounding boxes and the other one for penalizing the collision between the object meshes.

For the first constraint, we utilize the fact that based on the camera projection model, if we project predicted 3D bounding boxes onto the 2D image plane, the projected corners should be close to the 2D bounding box corners. This constraint therefore optimizes both camera pose and 3D bounding boxes. [28] imposes a similar loss where they penalize the deviation of the 2D projections of predicted 3D bounding box corners from ground truth 3D bounding box corners for both object bounding boxes and the room bounding box. However, since our model does not rely on any ground truth annotations in our described optimization process, we propose to use our detected 2D bounding boxes as a pseudo ground truth. We show the effectiveness of this loss term in Section 4. The formal definition of this term can be written as

$$\mathcal{L}_{\text{joint}}^j = \frac{1}{n_{\text{obj}}} \sum_{i=1}^{n_{\text{obj}}} \text{SmoothL}_2(f(X_i, c_i, \theta_i), h_i) \quad (2)$$

where c_i, θ_i are the size, centroid and orientation of the object i , h_i is the 2D bounding box estimate for object i , and f is a differentiable projection function that projects the corners of a 3D bounding box to a 2D image plane. Like [28], we use a smooth L_1 loss function comprised of a squared term if the absolute element-wise error falls below a threshold and an L_1 term otherwise.

Our second constraint is a loss term that penalize the collision between reconstructed object meshes. Although some pose estimation works [11] [18] have incorporated body collision losses, prior works in scene understanding have not explored this loss, because they either did not have the object shape information necessary to calculate the precise collision [18] [4], or did not take advantage of the object shape information that comes with the meshes [28]. We notice that inter-object collision is common in the output of human mesh estimation and was built upon [31] with the addition of scene contact (E_c) and penetration (E_p) terms by assuming access to ground truth scene scans. There are several differences between their full loss function and our

formulation in Fig. 1. First, we do not include any depth related terms, because we wish to perform estimation using solely RGB images whereas [11] propose model variants leveraging RGB-D information. Second, since we are performing joint estimation of the 3D scene from a monocular RGB image, we are not yet able to reason on scene contact or penetration after only human mesh estimation. So we include only a body self-penetration term in Fig. 1, which is computed following the approach in [1] [31] [41], and will consider human-scene constraints instead during our global optimization stage.

Within-scene losses. In Stage I of our approach, we also utilize within-scene constraints to generate an initial static scene estimation. Specifically, we design two within-scene constraints, one for encouraging 2D/3D consistency of the predicted object bounding boxes and the other one for penalizing the collision between the object meshes.

For the first constraint, we utilize the fact that based on the camera projection model, if we project predicted 3D bounding boxes onto the 2D image plane, the projected corners should be close to the 2D bounding box corners. This constraint therefore optimizes both camera pose and 3D bounding boxes. [28] imposes a similar loss where they penalize the deviation of the 2D projections of predicted 3D bounding box corners from ground truth 3D bounding box corners for both object bounding boxes and the room bounding box. However, since our model does not rely on any ground truth annotations in our described optimization process, we propose to use our detected 2D bounding boxes as a pseudo ground truth. We show the effectiveness of this loss term in Section 4. The formal definition of this term can be written as

$$\mathcal{L}_{\text{joint}}^j = \frac{1}{n_{\text{obj}}} \sum_{i=1}^{n_{\text{obj}}} \text{SmoothL}_2(f(X_i, c_i, \theta_i), h_i) \quad (2)$$

where c_i, θ_i are the size, centroid and orientation of the object i , h_i is the 2D bounding box estimate for object i , and f is a differentiable projection function that projects the corners of a 3D bounding box to a 2D image plane. Like [28], we use a smooth L_1 loss function comprised of a squared term if the absolute element-wise error falls below a threshold and an L_1 term otherwise.

Our second constraint is a loss term that penalize the collision between reconstructed object meshes. Although some pose estimation works [11] [18] have incorporated body collision losses, prior works in scene understanding have not explored this loss, because they either did not have the object shape information necessary to calculate the precise collision [18] [4], or did not take advantage of the object shape information that comes with the meshes [28]. We notice that inter-object collision is common in the output of human mesh estimation and was built upon [31] with the addition of scene contact (E_c) and penetration (E_p) terms by assuming access to ground truth scene scans. There are several differences between their full loss function and our

formulation in Fig. 1. First, we do not include any depth related terms, because we wish to perform estimation using solely RGB images whereas [11] propose model variants leveraging RGB-D information. Second, since we are performing joint estimation of the 3D scene from a monocular RGB image, we are not yet able to reason on scene contact or penetration after only human mesh estimation. So we include only a body self-penetration term in Fig. 1, which is computed following the approach in [1] [31] [41], and will consider human-scene constraints instead during our global optimization stage.

Within-scene losses. In Stage I of our approach, we also utilize within-scene constraints to generate an initial static scene estimation. Specifically, we design two within-scene constraints, one for encouraging 2D/3D consistency of the predicted object bounding boxes and the other one for penalizing the collision between the object meshes.

For the first constraint, we utilize the fact that based on the camera projection model, if we project predicted 3D bounding boxes onto the 2D image plane, the projected corners should be close to the 2D bounding box corners. This constraint therefore optimizes both camera pose and 3D bounding boxes. [28] imposes a similar loss where they penalize the deviation of the 2D projections of predicted 3D bounding box corners from ground truth 3D bounding box corners for both object bounding boxes and the room bounding box. However, since our model does not rely on any ground truth annotations in our described optimization process, we propose to use our detected 2D bounding boxes as a pseudo ground truth. We show the effectiveness of this loss term in Section 4. The formal definition of this term can be written as

$$\mathcal{L}_{\text{joint}}^j = \frac{1}{n_{\text{obj}}} \sum_{i=1}^{n_{\text{obj}}} \text{SmoothL}_2(f(X_i, c_i, \theta_i), h_i) \quad (2)$$

where c_i, θ_i are the size, centroid and orientation of the object i , h_i is the 2D bounding box estimate for object i , and f is a differentiable projection function that projects the corners of a 3D bounding box to a 2D image plane. Like [28], we use a smooth L_1 loss function comprised of a squared term if the absolute element-wise error falls below a threshold and an L_1 term otherwise.

Our second constraint is a loss term that penalize the collision between reconstructed object meshes. Although some pose estimation works [11] [18] have incorporated body collision losses, prior works in scene understanding have not explored this loss, because they either did not have the object shape information necessary to calculate the precise collision [18] [4], or did not take advantage of the object shape information that comes with the meshes [28]. We notice that inter-object collision is common in the output of human mesh estimation and was built upon [31] with the addition of scene contact (E_c) and penetration (E_p) terms by assuming access to ground truth scene scans. There are several differences between their full loss function and our

formulation in Fig. 1. First, we do not include any depth related terms, because we wish to perform estimation using solely RGB images whereas [11] propose model variants leveraging RGB-D information. Second, since we are performing joint estimation of the 3D scene from a monocular RGB image, we are not yet able to reason on scene contact or penetration after only human mesh estimation. So we include only a body self-penetration term in Fig. 1, which is computed following the approach in [1] [31] [41], and will consider human-scene constraints instead during our global optimization stage.

Within-scene losses. In Stage I of our approach, we also utilize within-scene constraints to generate an initial static scene estimation. Specifically, we design two within-scene constraints, one for encouraging 2D/3D consistency of the predicted object bounding boxes and the other one for penalizing the collision between the object meshes.

For the first constraint, we utilize the fact that based on the camera projection model, if we project predicted 3D bounding boxes onto the 2D image plane, the projected corners should be close to the 2D bounding box corners. This constraint therefore optimizes both camera pose and 3D bounding boxes. [28] imposes a similar loss where they penalize the deviation of the 2D projections of predicted 3D bounding box corners from ground truth 3D bounding box corners for both object bounding boxes and the room bounding box. However, since our model does not rely on any ground truth annotations in our described optimization process, we propose to use our detected 2D bounding boxes as a pseudo ground truth. We show the effectiveness of this loss term in Section 4. The formal definition of this term can be written as

$$\mathcal{L}_{\text{joint}}^j = \frac{1}{n_{\text{obj}}} \sum_{i=1}^{n_{\text{obj}}} \text{SmoothL}_2(f(X_i, c_i, \theta_i), h_i) \quad (2)$$

where c_i, θ_i are the size, centroid and orientation of the object i , h_i is the 2D bounding box estimate for object i , and f is a differentiable projection function that projects the corners of a 3D bounding box to a 2D image plane. Like [28], we use a smooth L_1 loss function comprised of a squared term if the absolute element-wise error falls below a threshold and an L_1 term otherwise.

Our second constraint is a loss term that penalize the collision between reconstructed object meshes. Although some pose estimation works [11] [18] have incorporated body collision losses, prior works in scene understanding have not explored this loss, because they either did not have the object shape information necessary to calculate the precise collision [18] [4], or did not take advantage of the object shape information that comes with the meshes [28]. We notice that inter-object collision is common in the output of human mesh estimation and was built upon [31] with the addition of scene contact (E_c) and penetration (E_p) terms by assuming access to ground truth scene scans. There are several differences between their full loss function and our

formulation in Fig. 1. First, we do not include any depth related terms, because we wish to perform estimation using solely RGB images whereas [11] propose model variants leveraging RGB-D information. Second, since we are performing joint estimation of the 3D scene from a monocular RGB image, we are not yet able to reason on scene contact or penetration after only human mesh estimation. So we include only a body self-penetration term in Fig. 1, which is computed following the approach in [1] [31] [41], and will consider human-scene constraints instead during our global optimization stage.

Within-scene losses. In Stage I of our approach, we also utilize within-scene constraints to generate an initial static scene estimation. Specifically, we design two within-scene constraints, one for encouraging 2D/3D consistency of the predicted object bounding boxes and the other one for penalizing the collision between the object meshes.

For the first constraint, we utilize the fact that based on the camera projection model, if we project predicted 3D bounding boxes onto the 2D image plane, the projected corners should be close to the 2D bounding box corners. This constraint therefore optimizes both camera pose and 3D bounding boxes. [28] imposes a similar loss where they penalize the deviation of the 2D projections of predicted 3D bounding box corners from ground truth 3D bounding box corners for both object bounding boxes and the room bounding box. However, since our model does not rely on any ground truth annotations in our described optimization process, we propose to use our detected 2D bounding boxes as a pseudo ground truth. We show the effectiveness of this loss term in Section 4. The formal definition of this term can be written as

$$\mathcal{L}_{\text{joint}}^j = \frac{1}{n_{\text{obj}}} \sum_{i=1}^{n_{\text{obj}}} \text{SmoothL}_2(f(X_i, c_i, \theta_i), h_i) \quad (2)$$

where c_i, θ_i are the size, centroid and orientation of the object i , h_i is the 2D bounding box estimate for object i , and f is a differentiable projection function that projects the corners of a 3D bounding box to a 2D image plane. Like [28], we use a smooth L_1 loss function comprised of a squared term if the absolute element-wise error falls below a threshold and an L_1 term otherwise.

Our second constraint is a loss term that penalize the collision between reconstructed object meshes. Although some pose estimation works [11] [18] have incorporated body collision losses, prior works in scene understanding have not explored this loss, because they either did not have the object shape information necessary to calculate the precise collision [18] [4], or did not take advantage of the object shape information that comes with the meshes [28]. We notice that inter-object collision is common in the output of human mesh estimation and was built upon [31] with the addition of scene contact (E_c) and penetration (E_p) terms by assuming access to ground truth scene scans. There are several differences between their full loss function and our

formulation in Fig. 1. First, we do not include any depth related terms, because we wish to perform estimation using solely RGB images whereas [11] propose model variants leveraging RGB-D information. Second, since we are performing joint estimation of the 3D scene from a monocular RGB image, we are not yet able to reason on scene contact or penetration after only human mesh estimation. So we include only a body self-penetration term in Fig. 1, which is computed following the approach in [1] [31] [41], and will consider human-scene constraints instead during our global optimization stage.

Within-scene losses. In Stage I of our approach, we also utilize within-scene constraints to generate an initial static scene estimation. Specifically, we design two within-scene constraints, one for encouraging 2D/3D consistency of the predicted object bounding boxes and the other one for penalizing the collision between the object meshes.

For the first constraint, we utilize the fact that based on the camera projection model, if we project predicted 3D bounding boxes onto the 2D image plane, the projected corners should be close to the 2D bounding box corners. This constraint therefore optimizes both camera pose and 3D bounding boxes. [28] imposes a similar loss where they penalize the deviation of the 2D projections of predicted 3D bounding box corners from ground truth 3D bounding box corners for both object bounding boxes and the room bounding box. However, since our model does not rely on any ground truth annotations in our described optimization process, we propose to use our detected 2D bounding boxes as a pseudo ground truth. We show the effectiveness of this loss term in Section 4. The formal definition of this term can be written as

$$\mathcal{L}_{\text{joint}}^j = \frac{1}{n_{\text{obj}}} \sum_{i=1}^{n_{\text{obj}}} \text{SmoothL}_2(f(X_i, c_i, \theta_i), h_i) \quad (2)$$

where c_i, θ_i are the size, centroid and orientation of the object i , h_i is the 2D bounding box estimate for object i , and f is a differentiable projection function that projects the corners of a 3D bounding box to a 2D image plane. Like [28], we use a smooth L_1 loss function comprised of a squared term if the absolute element-wise error falls below a threshold and an L_1 term otherwise.

Our second constraint is a loss term that penalize the collision between reconstructed object meshes. Although some pose estimation works [11] [18] have incorporated body collision losses, prior works in scene understanding have not explored this loss, because they either did not have the object shape information necessary to calculate the precise collision [18] [4], or did not take advantage of the object shape information that comes with the meshes [28]. We notice that inter-object collision is common in the output of human mesh estimation and was built upon [31] with the addition of scene contact (E_c) and penetration (E_p) terms by assuming access to ground truth scene scans. There are several differences between their full loss function and our

formulation in Fig. 1. First, we do not include any depth related terms, because we wish to perform estimation using solely RGB images whereas [11] propose model variants leveraging RGB-D information. Second, since we are performing joint estimation of the 3D scene from a monocular RGB image, we are not yet able to reason on scene contact or penetration after only human mesh estimation. So we include only a body self-penetration term in Fig. 1, which is computed following the approach in [1] [31] [41], and will consider human-scene constraints instead during our global optimization stage.

Within-scene losses. In Stage I of our approach, we also utilize within-scene constraints to generate an initial static scene estimation. Specifically, we design two within-scene constraints, one for encouraging 2D/3D consistency of the predicted object bounding boxes and the other one for penalizing the collision between the object meshes.

For the first constraint, we utilize the fact that based on the camera projection model, if we project predicted 3D bounding boxes onto the 2D image plane, the projected corners should be close to the 2D bounding box corners. This constraint therefore optimizes both camera pose and 3D bounding boxes. [28] imposes a similar loss where they penalize the deviation of the 2D projections of predicted 3D bounding box corners from ground truth 3D bounding box corners for both object bounding boxes and the room bounding box. However, since our model does not rely on any ground truth annotations in our described optimization process, we propose to use our detected 2D bounding boxes as a pseudo ground truth. We show the effectiveness of this loss term in Section 4. The formal definition of this term can be written as

$$\mathcal{L}_{\text{joint}}^j = \frac{1}{n_{\text{obj}}} \sum_{i=1}^{n_{\text{obj}}} \text{SmoothL}_2(f(X_i, c_i, \theta_i), h_i) \quad (2)$$

where c_i, θ_i are the size, centroid and orientation of the object i , h_i is the 2D bounding box estimate for object i , and f is a differentiable projection function that projects the corners of a 3D bounding box to a 2D image plane. Like [28], we use a smooth L_1 loss function comprised of a squared term if the absolute element-wise error falls below a threshold and an L_1 term otherwise.

Our second constraint is a loss term that penalize the collision between reconstructed object meshes. Although some pose estimation works [11] [18] have incorporated body collision losses, prior works in scene understanding have not explored this loss, because they either did not have the object shape information necessary to calculate the precise collision [18] [4], or did not take advantage of the object shape information that comes with the meshes [28]. We notice that inter-object collision is common in the output of human mesh estimation and was built upon [31] with the addition of scene contact (E_c) and penetration (E_p) terms by assuming access to ground truth scene scans. There are several differences between their full loss function and our

formulation in Fig. 1. First, we do not include any depth related terms, because we wish to perform estimation using solely RGB images whereas [11] propose model variants leveraging RGB-D information. Second, since we are performing joint estimation of the 3D scene from a monocular RGB image, we are not yet able to reason on scene contact or penetration after only human mesh estimation. So we include only a body self-penetration term in Fig. 1, which is computed following the approach in [1] [31] [41], and will consider human-scene constraints instead during our global optimization stage.

Within-scene losses. In Stage I of our approach, we also utilize within-scene constraints to generate an initial static scene estimation. Specifically, we design two within-scene constraints, one for encouraging 2D/3D consistency of the predicted object bounding boxes and the other one for penalizing the collision between the object meshes.

For the first constraint, we utilize the fact that based on the camera projection model, if we project predicted 3D bounding boxes onto the 2D image plane, the projected corners should be close to the 2D bounding box corners. This constraint therefore optimizes both camera pose and 3D bounding boxes. [28] imposes a similar loss where they penalize the deviation of the 2D projections of predicted 3D bounding box corners from ground truth 3D bounding box corners for both object bounding boxes and the room bounding box. However, since our model does not rely on any ground truth annotations in our described optimization process, we propose to use our detected 2D bounding boxes as a pseudo ground truth. We show the effectiveness of this loss term in Section 4. The formal definition of this term can be written as

$$\mathcal{L}_{\text{joint}}^j = \frac{1}{n_{\text{obj}}} \sum_{i=1}^{n_{\text{obj}}} \text{SmoothL}_2(f(X_i, c_i, \theta_i), h_i) \quad (2$$